

R- Final Exam

Shazman Khan

DD-MSBA-4

Q2

Code:

```
library(MASS)
```

```
library(rpart)
```

```
library(rpart.plot)
```

```
library(plotly)
```

```
library(ROCR)
```

```
summary(Aids2)
```

```
as.data.frame(Aids2)
```

```
aids=Aids2
```

```
2843*0.7
```

```
aids$status<-gsub(aids$status,pattern='A',replacement=1)
```

```
aids$status<-gsub(aids$status,pattern='D',replacement=0)
```

```
#Ans2a
```

```
train_index <- sample.split(aids, SplitRatio=0.7)
```

```
train <- aids[ train_index,]
```

```
test <- aids[!train_index,]
```

```
gc_train <- train
```

```
gc_test <- test
```

```
#Ans 2b
```

```
my_mod <- glm(age~death+T.categ+sex, data=gc_train, family=poisson)
```

```
my_mod
```

```
predict_logit <-predict(my_mod, gc_train, type='response')
```

```
predict_logit
```

Output:

```
Call: glm(formula = age ~ death + T.categ + sex, family = poisson,
data = gc_train)
```

Coefficients:

```
(Intercept)          death    T.categhsid      T.categid    T.categhet
T.categhaem  T.categblood T.categmother      -0.2750949    0.1483289
-0.2098921   -0.0858535   -2.6126071
T.categother          sexM
0.2780890    -0.1998361
```

Degrees of Freedom: 599 Total (i.e. Null); 590 Residual

Null Deviance: 1629

Residual Deviance: 1386 AIC: 4658

>

```
> predict_logit <- predict(my_mod, gc_train, type='response')
```

```
> predict_logit
```

```
261    177    416    195    886    437    983    5
05    164    302    902    135
38.386562 37.631460 38.209251 38.682399 36.672419 43.213116 36.672419 38.0160
46 37.244215 34.485919 37.520252 38.471016
669    969    568    838    388    795    694    2
89    74    919    915    308
36.933395 37.433498 37.073267 37.274786 38.221569 36.672419 37.301009 38.4450
97 38.816382 37.054804 28.281182 37.996000
932    356    21    999    548    254    459
```

Ans2c:

Designed the model by checking correlations and choosing variables that have a good correlations and then using it to predict values

From the regression model it is clear that age and T.Category are the most valuable variables

This tells us that the transmission of aids is largely dependent on age and the T.Categ which means that it is more susceptible for the youth while the category of transmission playing a huge role in if the patient survives or dies

.

Q3

Code:

Code

```
library(MASS)
```

```
library(rpart)
```

```
library(rpart.plot)
```

```
library(plotly)
```

```
library(ROCR)
```

```
summary(Aids2)
```

```
as.data.frame(Aids2)
```

```
aids=Aids2
```

```
2843*0.7
```

```
aids$status<-gsub(aids$status,pattern='A',replacement=1)
```

```
aids$status<-gsub(aids$status,pattern='D',replacement=0)
```

```
train_index <- sample.split(aids, SplitRatio=0.7)
```

```
train <- aids[ train_index,]
```

```
test <- aids[!train_index,]
```

```
gc_train <- train
```

```
gc_test <- test
```

```
#Ans3a
```

```
my_tree <- rpart(sex~age+T.categ+death, data=aids,method="class",control=rpart.control(cp=0.033))
```

```
rpart.plot(my_tree,type=1,extra=1, box.palette = c("pink","green"))
```

```
#Ans 3b
```

```
plotcp(my_tree)
```

```
predict_tree <- predict(my_tree, aids, type='prob')  
predict_tree
```

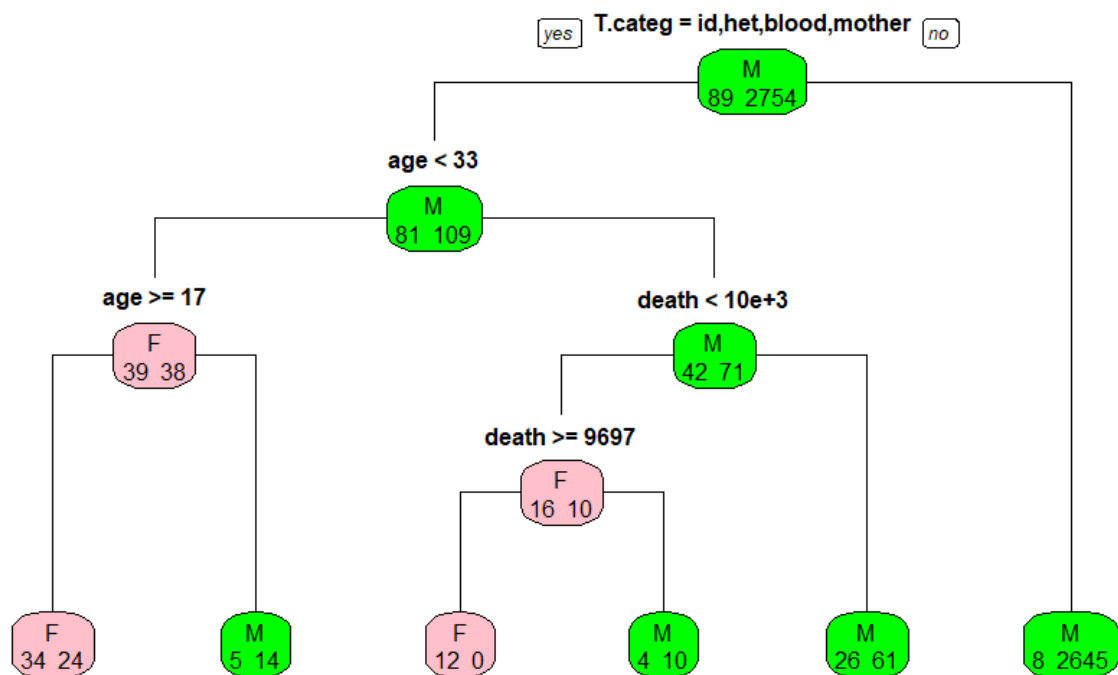

```
predict_logit <-predict(my_mod, gc_train, type='response')  
predict_logit
```

```
perf_tree<-performance(predict_tree)
```

```
perf_logit <- performance(predict_logit)
```

```
plot(perf_tree,col="black")
plot(perf_logit,col='blue',add=TRUE)
```

Output



Ans a: The tree

Ans B: As it can be seen from the relativity plot given by `plotcp(my_tree)` the optimal cp is at 0.033

Ans C: Its clear from the tree that age and death are the most valuable variables to predict if a patient lives or dies in the transmission tree. The biggest insight from the tree is that the ages of being above 17 for female is plays a vital role in terms of transmission and above the age of 15 for men.

For example if the age is lower than 42 for men the probability of death raises by $10e+3$

Ans D: In the performance graph it is clear to see that the tree performs better than the logistical regression and covers more area and is towards the left. Clearly the decision tree has performed better with this regard.

Q4

Ans 4a:

Output

Server code

```
#
# This is the server logic of a Shiny web application. You can run the
# application by clicking 'Run App' above.
#
# Find out more about building applications with Shiny here:
#
#   http://shiny.rstudio.com/
#
```

```
library(shiny)
library(rpart)
library(rpart.plot)
library(plotly)
# Define server logic required to draw a histogram
shinyServer(function(input, output) {
  output$distPlot <- renderPlot({
    library(readxl)
    mydf <- read.csv("german credit card.csv")

    mydf$purpose<-as.numeric(gsub("x","",mydf$purpose))

    mydf$good_bad<-gsub("good","1",mydf$good_bad)
    mydf$good_bad<-gsub("bad","0",mydf$good_bad)
    mydf$good_bad<-as.numeric(mydf$good_bad)

    ger_tree<-rpart(good_bad~age+amount+duration+checking,
                    data=mydf,method = "class",cp=input$bins)
    rpart.plot(ger_tree,extra=1, type = 1)
    t <- plot_ly(data = mydf, x =~age, y=~ good_bad)
    t
  })
  output$logit<- renderPrint({
    library(readxl)
    mydf <- read.csv("german credit card.csv")


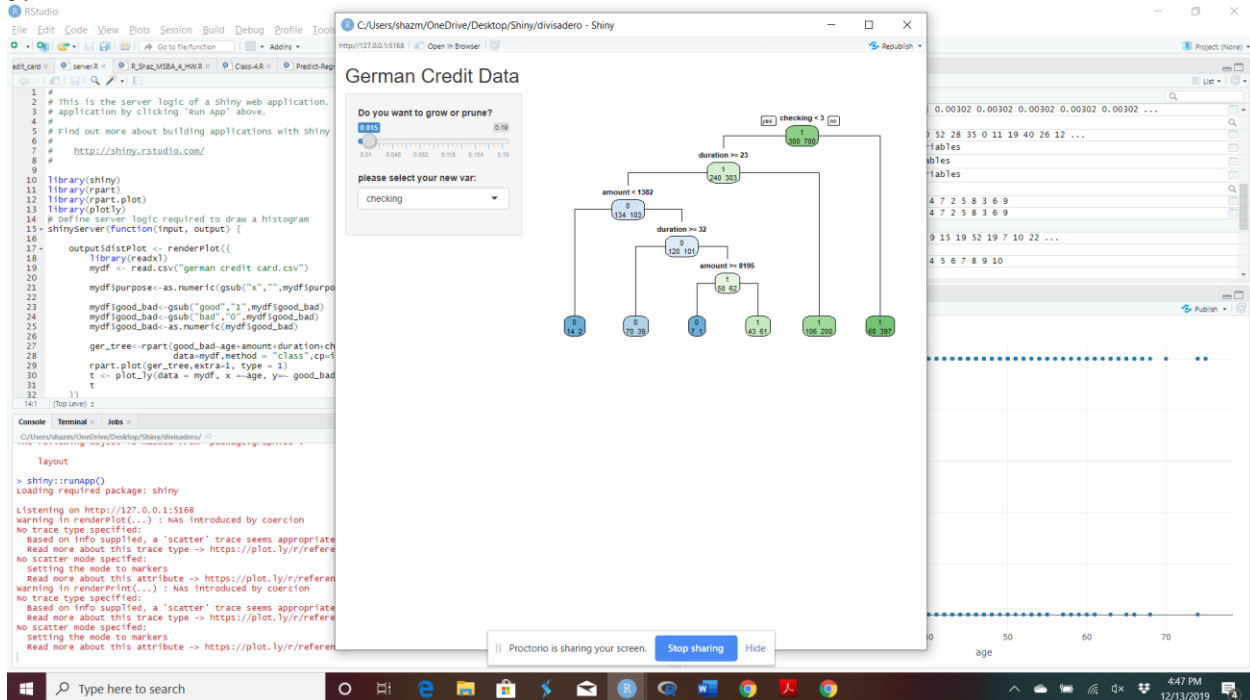
    mydf$purpose<-as.numeric(gsub("x","",mydf$purpose))

    mydf$good_bad<-gsub("good","1",mydf$good_bad)
    mydf$good_bad<-gsub("bad","0",mydf$good_bad)
  })
})
```

```

mydf$good_bad<-as.numeric(mydf$good_bad)
my_logit<-glm(paste("good_bad~age+amount+duration+",input$myvar),
              data=mydf,family = "binomial")
summary(my_logit)
t <- plot_ly(data = mydf, x =~age, y=~ good_bad)
t
})
})

```



Ans 4b:

Server Code#

This is the server logic of a Shiny web application. You can run the application by clicking 'Run App' above.

#

Find out more about building applications with shiny here:

#

<http://shiny.rstudio.com/>

#

```

library(shiny)

# Define server logic required to draw a histogram
shinyServer(function(input, output) {

  output$distPlot <- renderPlot({

    # generate bins based on input$bins from ui.R
    library(MASS)
    library(rpart)
    library(rpart.plot)
    library(plotly)
    library(ROCR)
    summary(Aids2)
    as.data.frame(Aids2)
    aids=Aids2
    2843*0.7
    aids$status<-gsub(aids$status,pattern='A',replacement=1)
    aids$status<-gsub(aids$status,pattern='D',replacement=0)

    train_index <- sample.split(aids, SplitRatio=0.7)
    train <- aids[ train_index,]
    test <- aids[!train_index,]
    gc_train <- train
    gc_test <- test
    my_mod <- glm(age~death+T.categ+sex, data=gc_train, family=poisson)
    my_mod

    predict_logit <-predict(my_mod, gc_train, type='response')
    predict_logit

    my_tree <- rpart(sex~age+T.categ+death, data=aids,method="class",cont
rol=rpart.control(cp=0.033))
    rpart.plot(my_tree,type=1,extra=1, box.palette = c("pink","green"))
    plotcp(my_tree)
  })

  output$logit<- renderPrint({
    library(MASS)
    library(rpart)
    library(rpart.plot)
    library(plotly)
    library(ROCR)
    summary(Aids2)
    as.data.frame(Aids2)
    aids=Aids2
    2843*0.7
    aids$status<-gsub(aids$status,pattern='A',replacement=1)
    aids$status<-gsub(aids$status,pattern='D',replacement=0)

    train_index <- sample.split(aids, SplitRatio=0.7)
    train <- aids[ train_index,]
    test <- aids[!train_index,]
    gc_train <- train
    gc_test <- test
    my_mod <- glm(age~death+T.categ+sex, data=gc_train, family=poisson)
    my_mod

    predict_logit <-predict(my_mod, gc_train, type='response')
    predict_logit

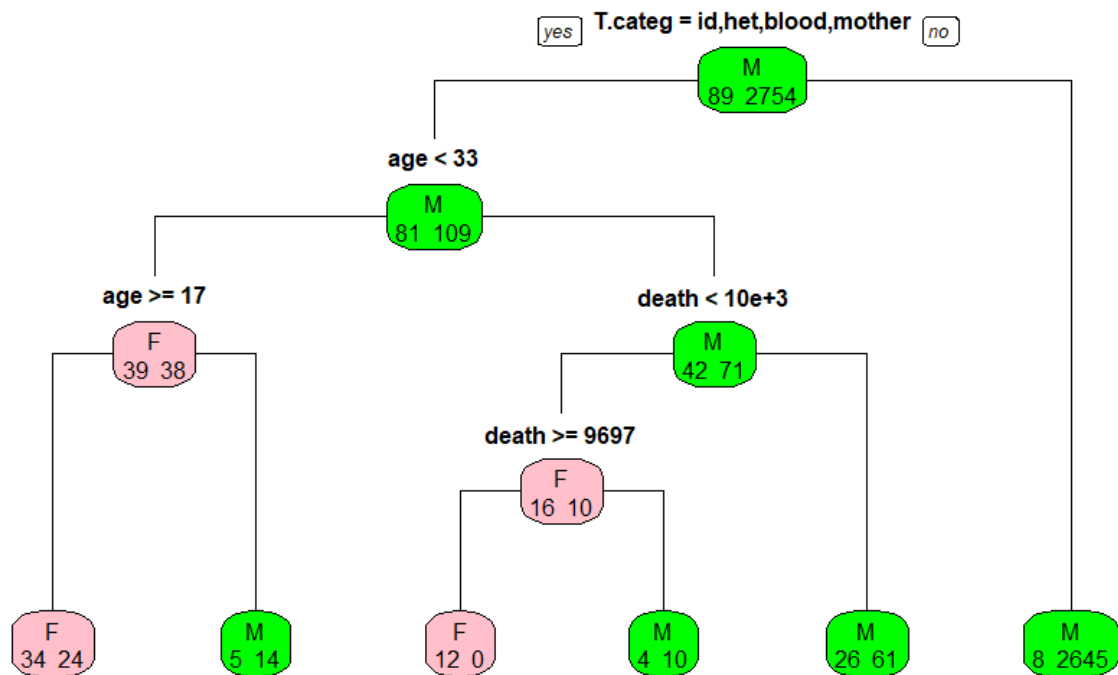
    my_tree <- rpart(sex~age+T.categ+death, data=aids,method="class",cont
rol=rpart.control(cp=0.033))
    rpart.plot(my_tree,type=1,extra=1, box.palette = c("pink","green"))
    plotcp(my_tree)
  })
})

```

```

})
})

```



Q5

Ans 5a

```

Code:
my_function <- function(net_income,total_assets,shareholders_equity,net_sales)
{
  ROA <- net_income/total_assets
  ROE <- net_income/shareholders_equity
  profit_margin <- net_income/net_sales
  ratios <- c(ROA,ROE,profit_margin)
  names <- c("ROA","ROE","Profit Margin")
  final <- cbind(ratios,names)
  print(final)
}
#Company 1
my_function(10000,40000,10000,50000)

#company 2
my_function(8000,45000,20000,35000)

#company 3
my_function(20000,60000,40000,70000)

#company 4
my_function(2000,10000,8000,4000)

```

Output
#company 1


```

my_function(10000,40000,10000,50000)
      ratios names
[1,] "0.25"  "ROA"
[2,] "1"     "ROE"
[3,] "0.2"   "Profit Margin"
>
> #company 2
> my_function(8000,45000,20000,35000)
      ratios names
[1,] "0.177777777777778" "ROA"
[2,] "0.4"               "ROE"
[3,] "0.228571428571429" "Profit Margin"
>
> #company 3
> my_function(20000,60000,40000,70000)
      ratios names
[1,] "0.333333333333333" "ROA"
[2,] "0.5"               "ROE"
[3,] "0.285714285714286" "Profit Margin"
>
> #company 4
> my_function(2000,10000,8000,4000)
      ratios names
[1,] "0.2"   "ROA"
[2,] "0.25"  "ROE"
[3,] "0.5"   "Profit Margin"

```

Q5B
ANS 5B

Code:

```

my_function <- function(net_income,total_assets,shareholders_equity,net_sales)
{
  ROA <- net_income/total_assets
  ROE <- net_income/shareholders_equity
  profit_margin <- net_income/net_sales
  ratios <- c(ROA,ROE,profit_margin)
  names <- c("ROA","ROE","Profit Margin")
  final <- cbind(ratios,names)
  print(final)
}
#Company 1
my_function(10000,40000,10000,50000)

#company 2
my_function(8000,45000,20000,35000)

#company 3
my_function(20000,60000,40000,70000)

#company 4
my_function(2000,10000,8000,4000)

Avg_roa= (0.177+0.333+0.2+0.25)/4
Avg_roa

#company weighted roa
ROA1=Avg_roa*0.25
ROA1
ROA2=Avg_roa*0.177
ROA2

```

```
ROA3=Avg_roa*0.333
ROA3
ROA4=Avg_roa*0.2
ROA4
```

Output

```
ROA1=Avg_roa*0.25
> ROA1
[1] 0.06
> ROA2=Avg_roa*0.177
> ROA2
[1] 0.04248
> ROA3=Avg_roa*0.333
> ROA3
[1] 0.07992
> ROA4=Avg_roa*0.2
> ROA4
[1] 0.048
```

Q5d

Ans 5d

The ROA Tells us if the company has optimized its assets to drive value. The higher the ROA the better the company's performance. Whether a company is doing well or not depends on the ROA vs the industry standard since this isn't known now it's not easy to tell how the company is faring in the industry. The Return on equity tells us if we are going to make our money back if we invest in the company. A company would strive to achieve a better return on equity. But since the industry standard is not known it is difficult to figure out which firm to invest in based on ROE. Profit margin tells us the profit the firm has earned vs costs. The higher the ratio the better. In conclusion it is easy to see that in between these 4 companies company 3 has been doing better hence I would invest in company 3.

Q5E

Ans 5e

Code

```
#Investing 10% from company 1 into company 4
my_function(3000,14000,9000,9000)
```

Output:

```
      ratios      names
[1,] "0.214285714285714" "ROA"
[2,] "0.333333333333333" "ROE"
[3,] "0.333333333333333" "Profit Margin"
```

>

As we can see that ROA AND ROE of company 4 increases from such an investment but its profit margin drops quite a bit.

Q1

Ans1a)

Code:

```
barto <- rexp(1000,rate=0.3)
barto
```

Output:

```
[1] 1.514128384 2.442114115 0.086278031 4.660980403 0.667792610 2.06980
5072 0.719387551 1.876248484
[9] 3.825787871 0.225954320 6.657387037 0.659825224 2.336584789 2.91
9798481 1.070630719 3.464839956
[17] 1.056264696 1.679832808 4.522413841 0.334369801 4.009987588 1.23
1910288 2.313090206 0.503832835
[25] 2.199534969 3.215783255 2.736417966 4.107485162 1.414347352 0.33
7204933 2.284926310 0.509329541
[33] 7.571780920 1.619301016 4.497634098 1.064059289 0.728062259 3.17
4738840 1.379020219 2.733223089
[41] 3.438464505 4.843509075 1.786017818 17.902562928 0.151506290 7.45
1834150 2.888688594 7.445912230
[49] 11.109537829 6.221050216 10.643293907 6.230059527 1.486908801 0.91
9346235 1.002389531 0.092402954
[57] 1.096363697 0.915655123 3.586564784 0.203761771 1.801614358 1.67
3699513 1.786267599 8.088963366
[65] 12.064868669 3.856239486 2.262524120 2.910845920 4.808643609 2.39
2932810 9.608325664 1.868552306
```

Ans1b:

Code:

```
barto <- rexp(1000,rate=0.3)
barto
mean(barto)
var(barto)
```

Output:

```
mean(barto)
[1] 3.218099
> var(barto)
[1] 9.969917
```

Ans1c

Code:

```
a)
exp(0.3-barto*3)
mean(barto)
```

Output:

```
1] 1.437327e-02 8.881745e-04 1.042025e+00 1.141758e-06 1.820675e-01 2.713773e
-03 1.559589e-01 4.850151e-03
[9] 1.398726e-05 6.853244e-01 2.860810e-09 1.864717e-01 1.218964e-03 2.119
001e-04 5.437275e-02 4.130529e-05
[17] 5.676734e-02 8.743030e-03 1.730255e-06 4.950436e-01 8.048999e-06 3.351
605e-02 1.307981e-03 2.977508e-01
[25] 1.838868e-03 8.719620e-05 3.673285e-04 6.007779e-06 1.938914e-02 4.908
509e-01 1.423297e-03 2.928811e-01
```